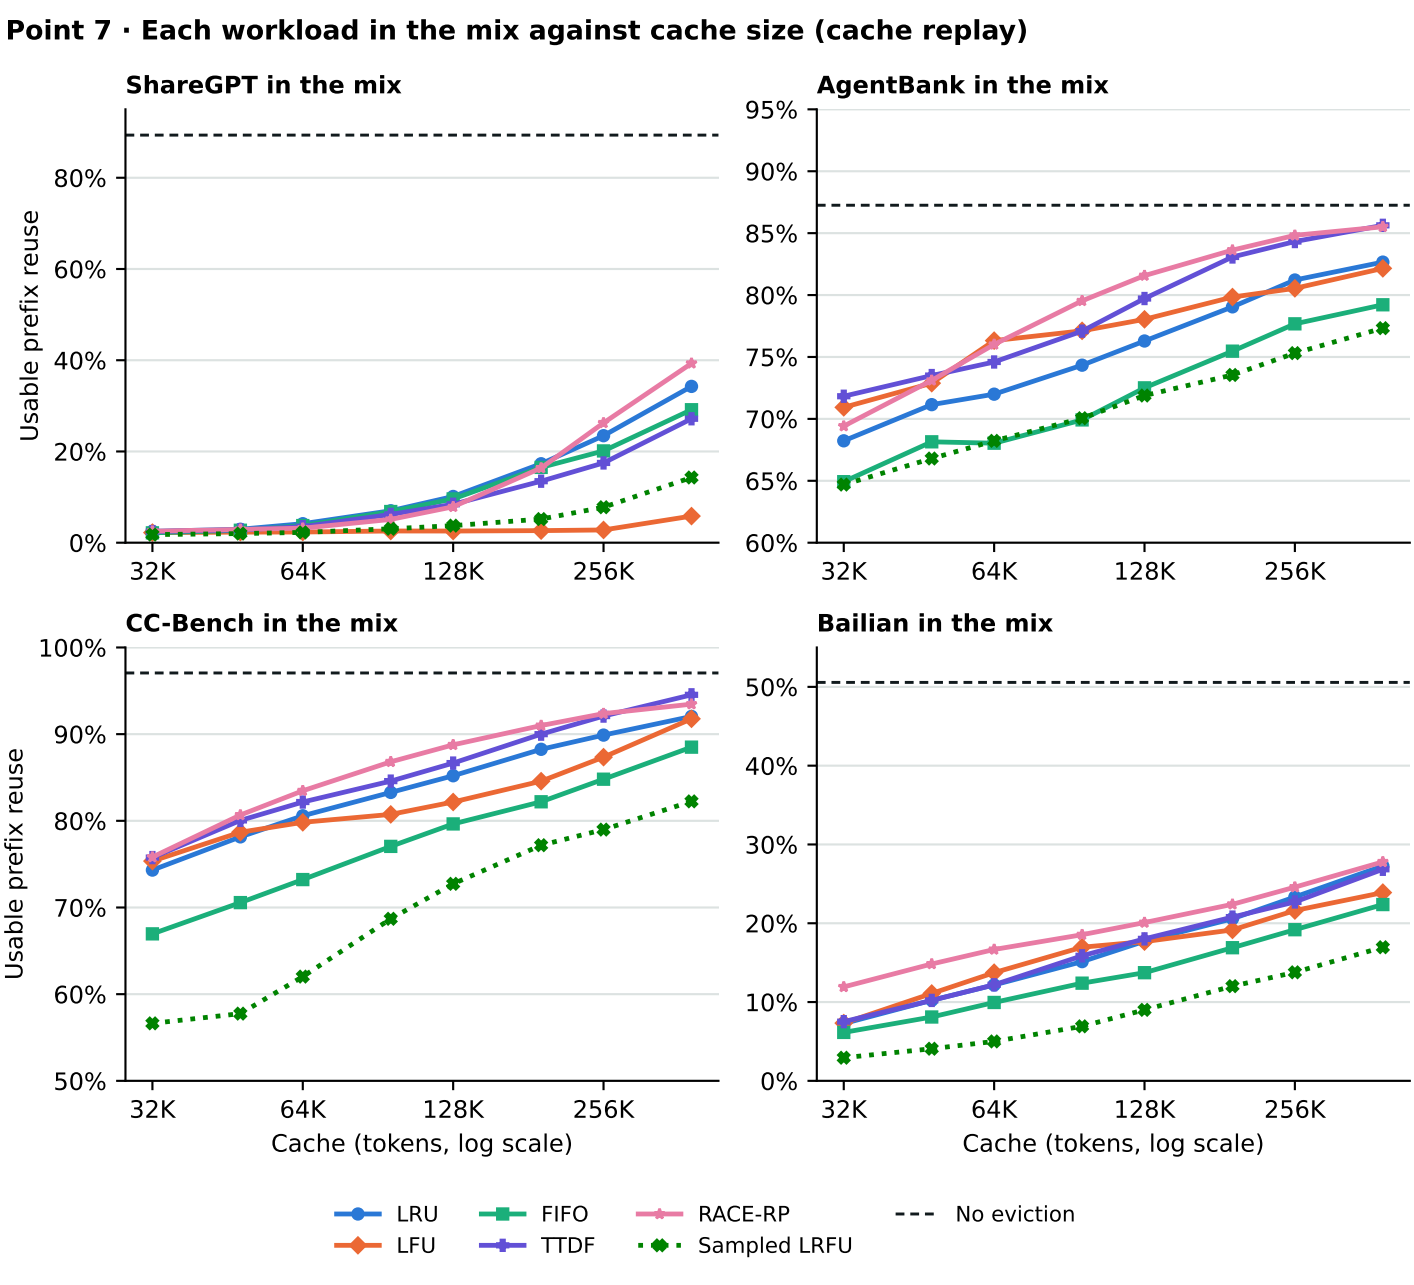
Point 7 · Each workload in the mix against cache size (cache replay)
ShareGPT in the mix
AgentBank in the mix
CC-Bench in the mix
Bailian in the mix
Usable prefix reuse
Cache (tokens, log scale)
0%
20%
40%
60%
80%
50%
65%
70%
75%
85%
90%
95%
100%
10%
30%
32K
64K
128K
256K
LRU
LFU
FIFO
TTDF
RACE-RP
Sampled LRFU
No eviction